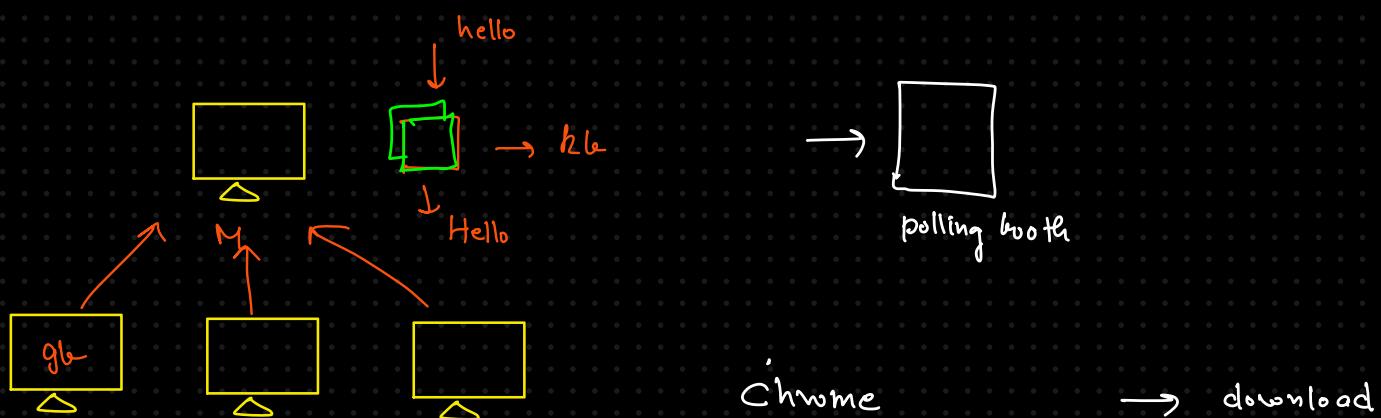


Map Reduce

data → store ✓
process ✓

Code locality



data is moving towards the code



data locality

+

Concept of local processing

Send code to
data

We are going to aim for distributed processing → Map Reduce

Framework

Map Reduce

GMR \rightarrow MR

Map reduce is a programming model designed for processing large dataset in a distributed computing environment.

It simplifies handling massive data by

dividing task into smaller, manageable subtasks and distributing them across multiple machine in a cluster.

e.g. → vote count
word count
product sales

→ write understanding distributed processing.

→ Spark \Rightarrow 70%

A 100	500	700

B 50	600	300

1. We are going to have data in a distributed manner via hdfs.

→ Map

→ Reduce

input data \rightarrow split into chunk



mapper



intermediate output



reducer



final output

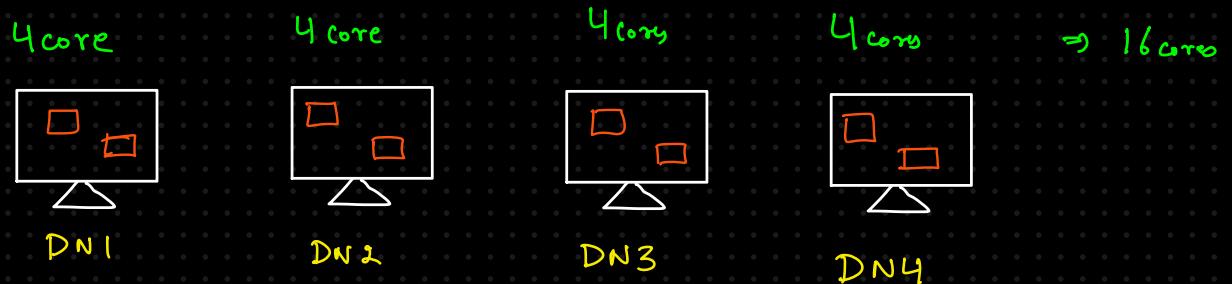
Traditional lang code like we have been writing doesn't work in a distributed way.

Map Reduce and Cluster

400 mb

4 node cluster

DN CPU + RAM + HDD



1 cpu core can process 1 mapper job on 1 block / input split of data

32 blocks → we can process 16 core via above cluster ⇒ we will not be failing the processing rather it is going to happen in batches & will take time

if we decrease a block size ⇒ parallelism ↑↑

2 blocks ← 512 mb ← 1gb → 128 mb block size = 8 blocks
└ 64 mb ⇒ 16 blocks
└ 4 mb ⇒ 64 blocks

1. Distributed Processing

2. Fault tolerant

3. Data Locality

4. Scalability

Map Reduce - Word Count example

We want to count the occurrence of each word in a file via the map reduce framework.

1. Input data (input.txt)

2. Mapper

no. of mapper task = no. of blocks

(1, gohu vegeta gohan)

(K, V) record reader converts the input to (K, V) pair for mapper input

Mapper input \rightarrow (1, gohu vegeta gohan)

\downarrow (gohu, 1)

(vegeta, 1)

(gohan, 1)

(Gohu, 2), (Vegeta, 2)

\uparrow Mappers output

3. Combiner (Optional)

without combiner: mapper sends all the key value pairs to the reducers, increasing shuffle data

with combiner: Partial aggregation at each mapper to reduce shuffle data.

((Gohu, 1), (Gohu, 1)) \longrightarrow (Gohu, 2)

4. Shuffle and sort phase

\rightarrow data from all mapper is shuffled and grouped by Key

Gohu [1, 2, 2, 1]

Gohan [1, 1, 1]

\rightarrow Keys are sorted for efficient reduction

5. Reducers

Each reducer processes the grouped Key value pair and aggregates value.
we can control the numbers of reducer

6. Final Output

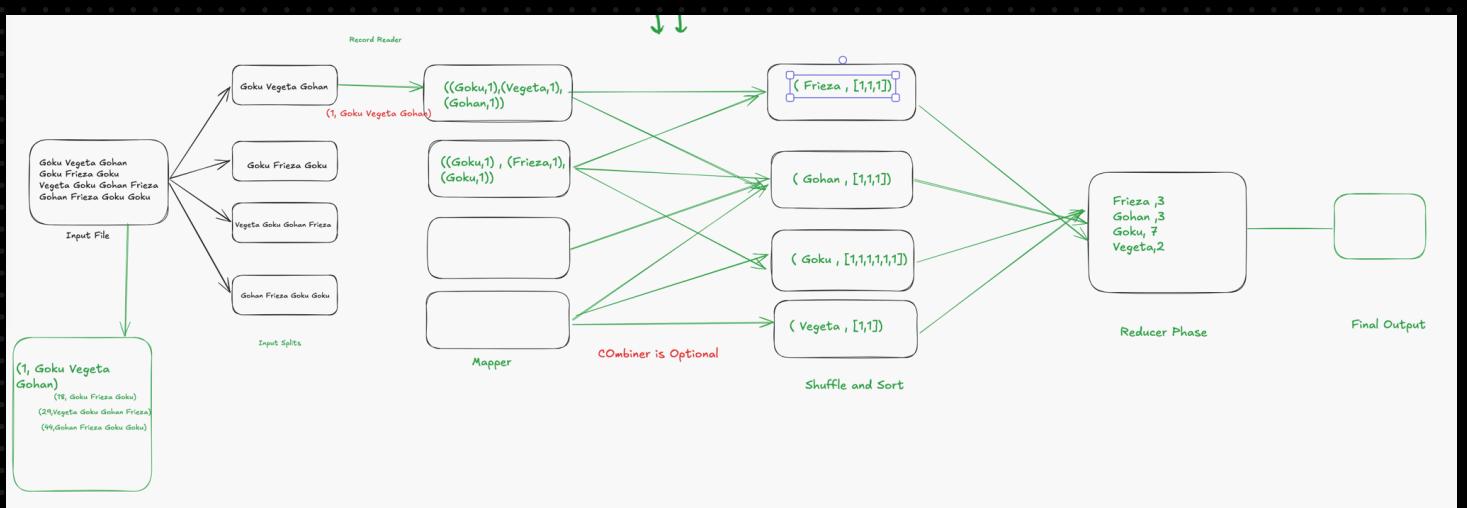
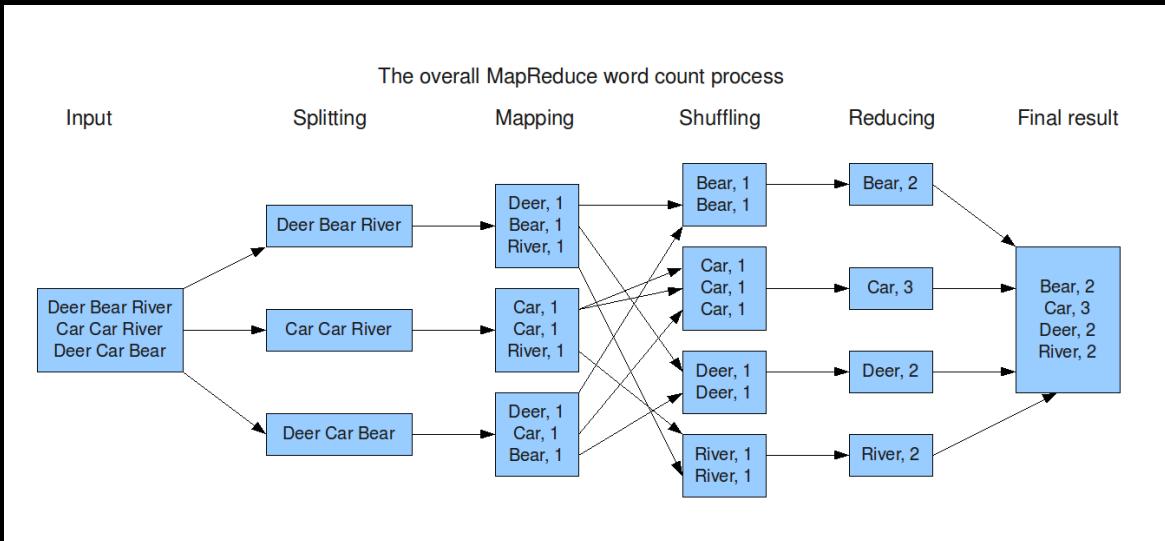
final output is written to hdfs or any other location.

Reducer Input [Goku, [1, 1, 1]]



Reducer output [Goku, 3]

Step	Process	Output Example
Input Splitting	Splitting input file into HDFS chunks.	Split 1: Goku Vegeta Gohan , Split 2: Goku Frieza Goku , etc.
Mapper Phase	Generating key-value pairs (word, 1) from each split.	("Goku", 1), ("Vegeta", 1), ("Gohan", 1)
Combiner Phase	Optional local aggregation at Mapper to reduce shuffle data.	("Goku", 2), ("Frieza", 1)
Shuffle and Sort	Grouping and sorting intermediate key-value pairs by key.	("Goku", [1, 2, 1, 2]), ("Frieza", [1, 1, 1])
Reducer Phase	Aggregating grouped data to produce final results.	("Goku", 6), ("Frieza", 3)
Final Output	Writing final aggregated results to HDFS.	Frieza 3 , Gohan 3 , Goku 6 , Vegeta 2



Map Reduce - Practical on Hadoop Cluster

⇒ Normal Map Reduce with 1 reducer



Step 2:

2 important things

1. Of course MR is optimized for large inputs.
 2. A single reducer can get bottlenecked or overloaded by lots of data.

Map Reduce Practical : More than one Reducer



We can have multiple reducers to split the workload , parallel processing of grouped keys is achieved via this

1. We have to make sure that we do / employ mappers to achieve parallelism & do maximum heavy lifting

inc.
10 mapper → 10 min
1 reducer → 40 min
↓
2 reducers ⇒ 5 min
50 min → 15 min

2. We cannot have chain of reducers layers
We can have a single reducer layer

Partitioning

determine how intermediate keys are distributed to reducers.

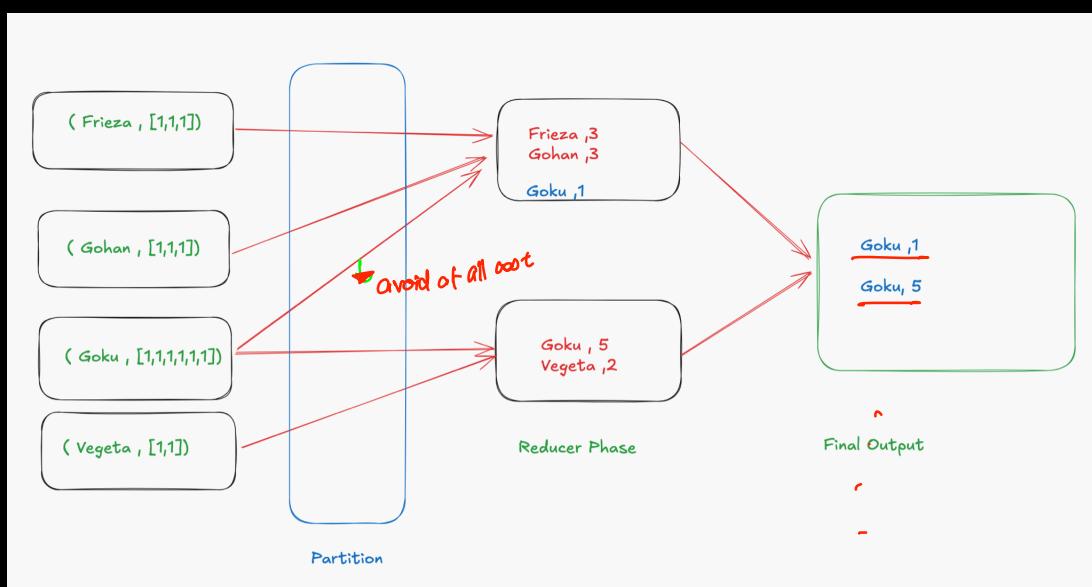
ensure that same key goes to same reducer

hashing

$$a-m \rightarrow 1$$
$$n-z \rightarrow 2$$

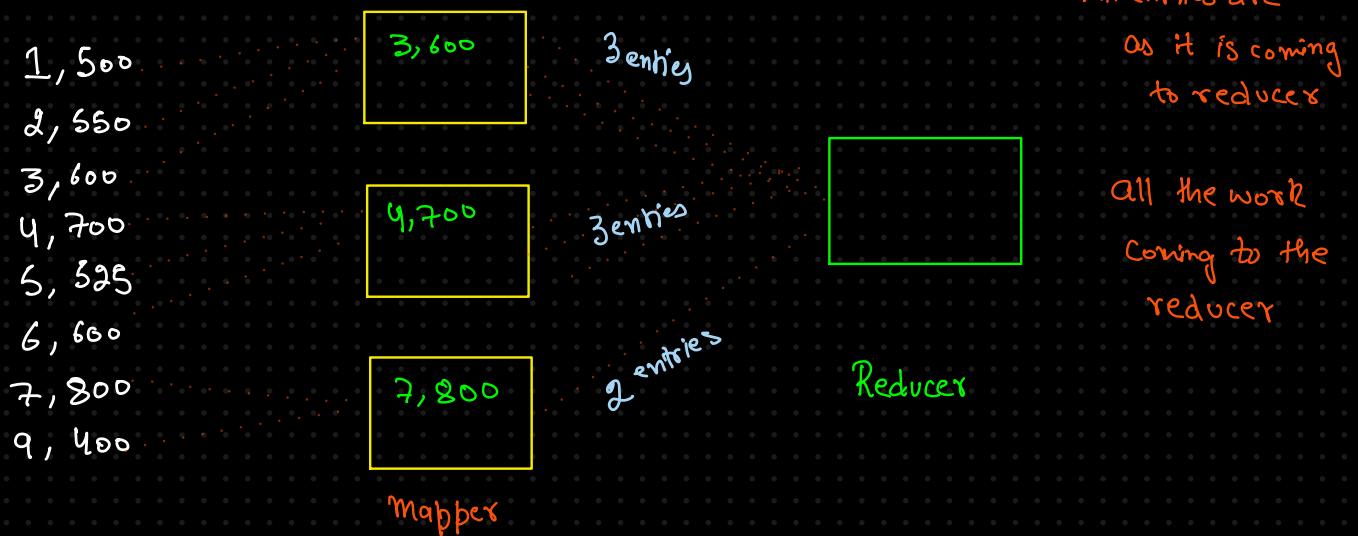
modulus operator

%



Combiner

⇒ A mini-reducer that performs partial aggregation at mapper level.



Using Combiner ↴

We should have calculated local answer in each of the mappers.
& then our reducer would just have to check for 3 entries.

It reduces the amount of data getting moved over the network

Q: Can we always use Combiner?

Find the median, average

So for non-associative or non-commutative operation we cannot use Combiner.

main priority is getting right

Map Reduce with Zero Reducers

- when the output of mapper layer is the final output
- filtering, data transformation or format conversion.

Goku vegeta friega $\xrightarrow{>4}$ Vegeta friega

When we don't require any aggregation.

Map Reduce on a log file + Industry Use case

Log file

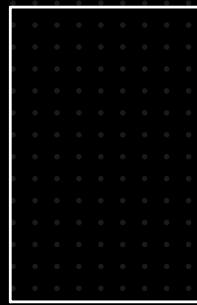
filtering , group by

hive . \Rightarrow abstraction of MR
Pig Latin

Amazon

{ iPhone
Samsung
Huawei
Nokia
TV

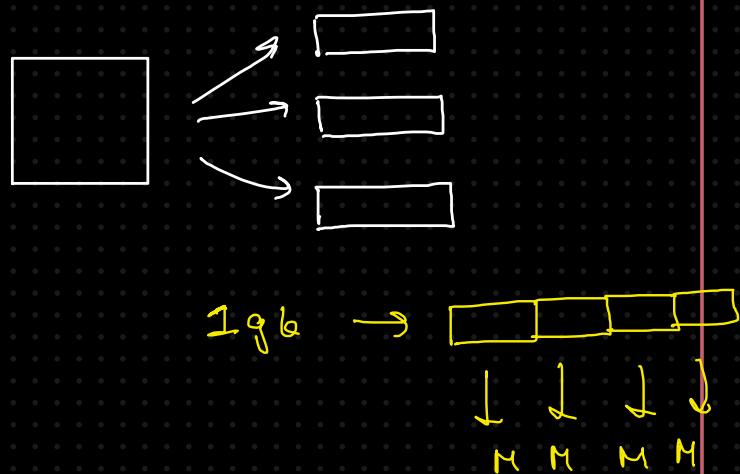
Phone, Mobile, Apple
Handheld phone



Input Split

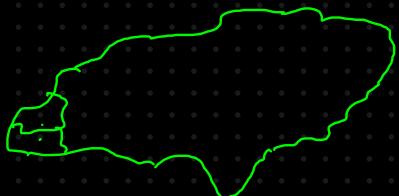
block / input split

An Input split is a logical division of the input data used by Hadoop MR to process large dataset efficiently in a distributed manner.



The core can run one mapper job on one split of data.

e.g. ⇒ countries on earth / land

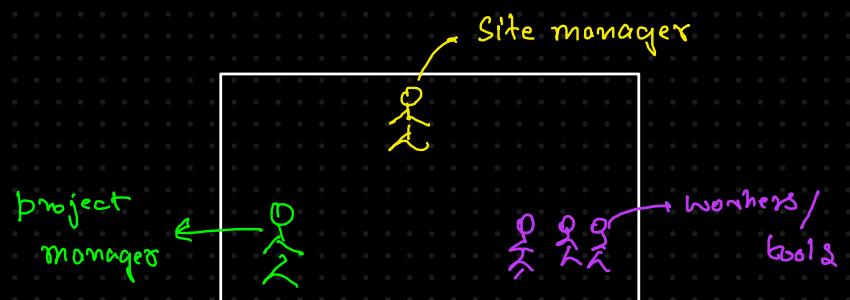


YARN analogy

Construction

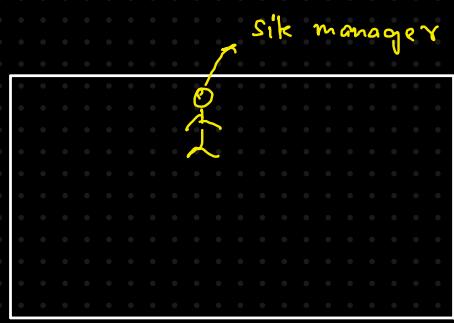


Central office



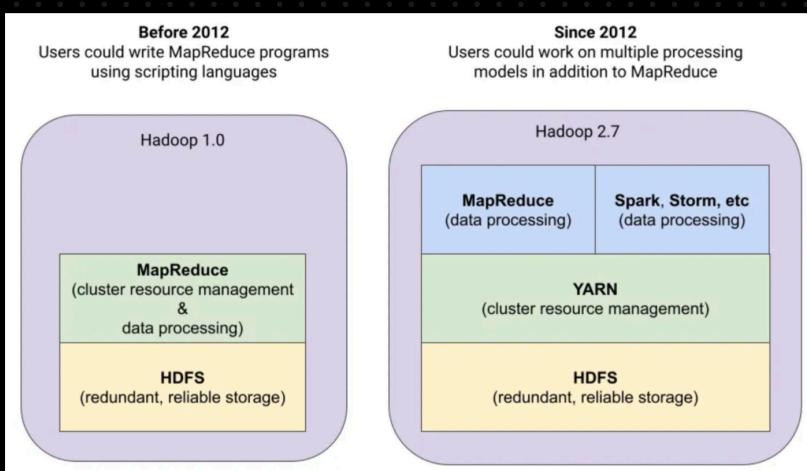
Construction site

Interview



Construction Site

YARN (Yet Another Resource Negotiator)



Interview

Spark

still getting used

↳ dataproc cluster

resource management layer in hadoop.

components of Yarn

1. Resource Manager: master daemon of YARN

Central office

Allocate resources

Efficiently resource usage across jobs

(a) Scheduler

(b) Application Master: manages application in our cluster

Application master

2. Node Manager: uploads everything to Resource manager

Site manager

keeps the data

to RM updated

Container

3. Application Master: negotiates resources with RM

≈ no. of jobs

works with node manager

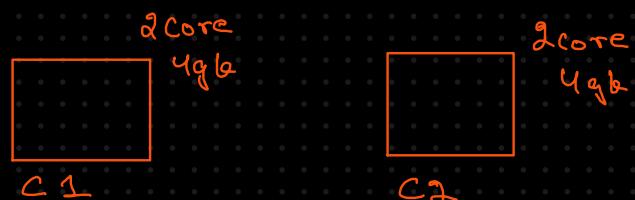
for executing and monitoring the task

Project manager

4. Containers : physical resource
(RAM, core, disk)

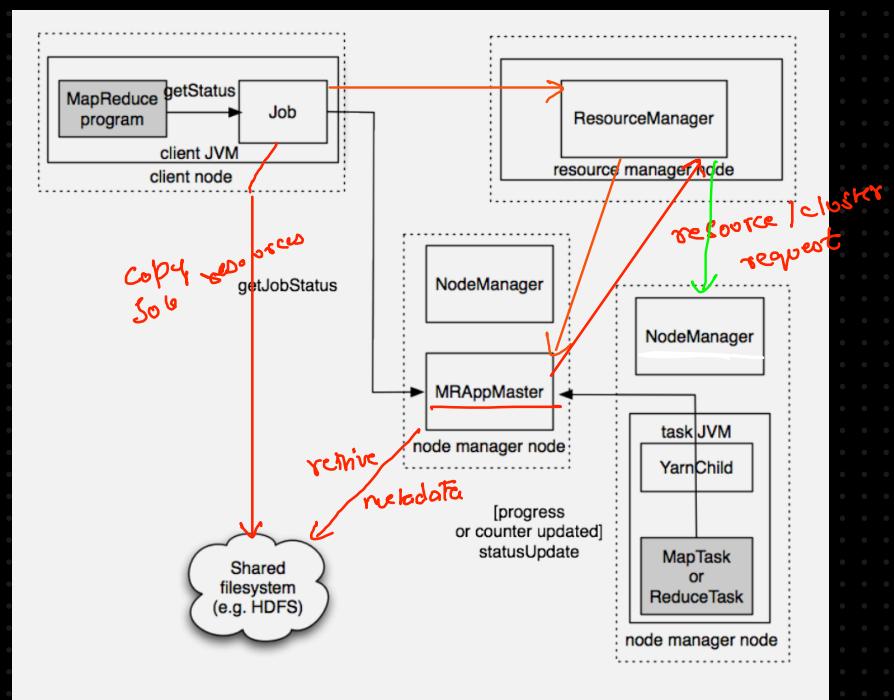
team of workers
\$ tools

executes the task assigned by the Node Manager



YARN Process - Step by Step process

1. Job submission
2. Launching of Application master
3. Requesting Containers
4. Allocating Resources (RM → NM)
5. Executing tasks
6. Monitoring Progress
7. Job completion



RM needs to free up the resources for future jobs.

